

Universitat Politècnica de València

DSIC - Departament de Sistemes Informàtics i Computació

Academic Project

MIARFID - Master's in Artificial Intelligence and Pattern
Recognition

Optimization of Keyboard Layouts for English

Author:

Jordi Cantavella Ferrero

Academic Year 2025-2026

Contents

1	Introduction	5
1.1	Context	5
1.2	Objectives	5
1.3	Scope and Limitations	6
2	Problem Description and Theoretical Framework	7
2.1	History of Keyboard Layouts	7
2.2	Problem Modeling	8
2.2.1	Individual Representation	8
2.2.2	Fitness Function	8
2.3	Linguistic Analysis and Datasets	9
3	Genetic Algorithm: Methodology and Experiments	11
3.1	Genetic Algorithm Implementation	11
3.1.1	Initial Population	11
3.1.2	Selection: Tournament Selection	11
3.1.3	Crossover: Two-Point Crossover	12
3.1.4	Mutation: Swap Mutation	12
3.1.5	Replacement: Elitist Strategy	12
3.2	Experimental Design	12
3.2.1	Experiment Series	13
3.3	Experimental Results and Analysis	13
3.3.1	Experiment 1: Population Size	13
3.3.2	Experiment 2: Tournament Selection Size	14
3.3.3	Experiment 3: Mutation Rate	15
3.3.4	Experiment 4: Elite Percentage	16
3.3.5	Summary of Parameter Experiments	17
3.4	Simulated Annealing	19

4 Conclusions	20
4.1 Summary of the Developed Work	20
4.2 Achieved Objectives	20
4.3 Future Work	21
A Detailed Experimental Results	22
A.1 Experiment 1: Population Size - Detailed Results	22
A.2 Experiment 2: Tournament Selection - Detailed Results	23
A.3 Experiment 3: Mutation Rate - Detailed Results	23
A.4 Experiment 4: Elite Percentage - Detailed Results	24
B Algorithm Implementation Details	25
B.1 Computational Optimization Strategies	25
B.2 Parameter Sensitivity Summary	26
B.3 Convergence Pattern Analysis	26
C Comparison with Existing Layouts	28
C.1 Detailed Layout Comparisons	28
C.2 Cross-Corpus Generalization Analysis	29
C.3 Letter Placement Comparison	29

List of Figures

2.1	Comparison of the QWERTY, Dvorak, and Colemak keyboard layouts.	7
2.2	Finger mapping visualization showing key-to-finger assignments.	8
3.1	Experiment 1: Impact of Population Size on convergence	14
3.2	Experiment 2: Impact of Tournament Selection Size	15
3.3	Experiment 3: Impact of Mutation Rate	16
3.4	Experiment 4: Impact of Elite Percentage	17

List of Tables

3.1	Baseline Configuration for Parameter Experiments	12
3.2	Optimal Parameter Configuration	18
A.1	Final fitness values for different population sizes	22
A.2	Final fitness values for different tournament sizes	23
A.3	Final fitness values for different mutation rates	23
A.4	Final fitness values for different elite percentages	24
B.1	Parameter Sensitivity Analysis	26
C.1	Comprehensive Layout Comparison - Moby Dick Corpus . . .	28
C.2	Comprehensive Layout Comparison - Wizard of Oz Corpus .	28
C.3	Generalization Performance: Layouts Tested on Alternative Corpus	29
C.4	Home Row Letter Distribution Across Layouts	29

Chapter 1

Introduction

1.1 Context

The layout of current keyboards is not optimized for ergonomics and typing efficiency. Current layouts have their origins in traditional typewriting and have not evolved to meet the needs of current users. Layouts like *QWERTY*, which emerged in 1873, were designed for mechanical typewriters and are no longer the most efficient for typing on digital devices. Studies showed that the *QWERTY* layout generated inefficient movements, contributing to fatigue and the risk of injury. This led to subsequent layouts such as *Dvorak* (1936) and *Colemak* (2006), which aimed to improve efficiency and reduce user fatigue. Trying to maximize typing speed and minimize physical effort, these layouts are based on ergonomic principles and linguistic analysis.

1.2 Objectives

The main objective of this work is the optimization of keyboard layouts, aiming to generate ergonomic and efficient keyboard arrangements. To achieve this, a genetic algorithm and simulated annealing will be implemented to explore the space of possible layouts and find optimal configurations based on ergonomic and linguistic criteria. The following specific objectives are defined:

- Implement a genetic algorithm for generating and evaluating different keyboard layouts.

- Implement a simulated annealing algorithm for the same problem.
- Define evaluation metrics that consider distance, hand alternation, and finger effort.
- Systematically analyze the impact of algorithm parameters on optimization performance.
- Compare the generated layouts with existing ones such as *QWERTY*, *Dvorak*, and *Colemak*.

1.3 Scope and Limitations

Given the time and resource constraints, this work will focus on optimizing keyboard layouts for the English language. Aspects related to the physical implementation of keyboards or adaptation to different devices will not be addressed.

Also given the different characteristics of each distribution, the optimization will be focused on the most common Latin alphabet layouts, including all the letters, and the most common punctuation marks.

Chapter 2

Problem Description and Theoretical Framework

2.1 History of Keyboard Layouts

The **QWERTY** keyboard, designed in 1873 by Christopher Latham Sholes, was created to prevent the keys of early mechanical typewriters from jamming by separating frequently used letter pairs. The **Dvorak** keyboard (1936) was scientifically designed to improve efficiency, reducing finger movement by up to 35% compared to QWERTY. The **Colemak** keyboard (2006) represents a modern evolution that balances efficiency with ease of transition from QWERTY.

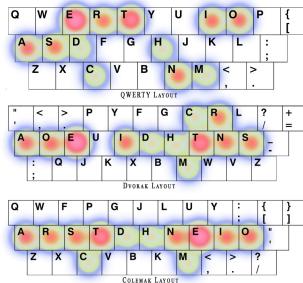


Figure 2.1: Comparison of the QWERTY, Dvorak, and Colemak keyboard layouts.

2.2 Problem Modeling

The keyboard layout optimization problem is modeled as a combinatorial optimization problem where the goal is to find the best arrangement of 30 characters (26 letters plus 4 punctuation marks: period, comma, semicolon, apostrophe) [a..z,,;,'], representing the genotype.

The phenotype is represented as a permutation of these characters arranged on a 3x10 grid to minimize typing cost. Easing the visualization of the keyboard layout and key positions.

2.2.1 Individual Representation

Each keyboard layout is represented as a permutation of 30 characters arranged in a standard grid structure:

0, 1, 2, 3, 4, 5, 6, 7, 8, 9	(Top Row)
10, 11, 12, 13, 14, 15, 16, 17, 18, 19	(Home Row)
20, 21, 22, 23, 24, 25, 26, 27, 28, 29	(Bottom Row)

Each position is assigned to a specific finger based on standard touch typing conventions. Figure 2.2 illustrates the finger assignments across the keyboard grid.

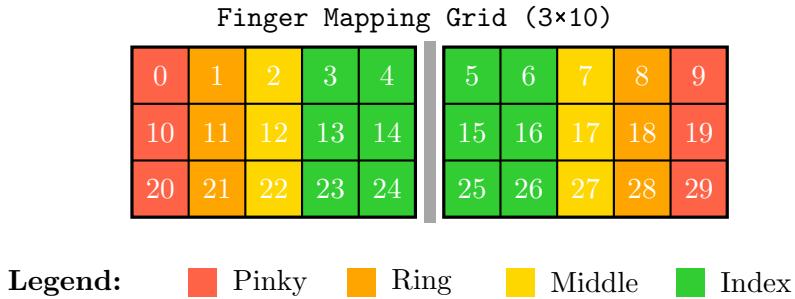


Figure 2.2: Finger mapping visualization showing key-to-finger assignments.

2.2.2 Fitness Function

The fitness function evaluates typing cost by analyzing consecutive character pairs (bigrams) in the text. For each bigram, the cost is computed as:

$$\text{cost}(c_i, c_{i+1}) = d_{\text{euclidean}}(c_i, c_{i+1}) \times \max(1.0 + P_{\text{finger}}(c_i, c_{i+1}), 0.1) \quad (2.1)$$

where $d_{\text{euclidean}}$ is the Euclidean distance between key positions and P_{finger} is a penalty system based on the following ergonomic factors:

- **Same Finger Usage:** Heavy penalty (1.0) when consecutive characters use the same finger, with additional penalty (2.0) for weak fingers (pinky, ring).
- **Same Hand Usage:** Moderate penalty (1.0) to promote hand alternation.
- **Row Transitions:** Penalties for vertical movement (0.2 for adjacent rows, 0.8 for two-row jumps), with additional costs for weak fingers.
- **Weak Finger Usage:** Penalties for pinky (0.15) and ring finger (0.1).
- **Same Column Movement:** Extra penalty (0.3) for vertical movement, with additional costs for outer columns.

The total fitness is the sum of all bigram costs:

$$F(\mathcal{L}) = \sum_{i=1}^{n-1} \text{cost}(c_i, c_{i+1}) \quad (2.2)$$

The objective is to **minimize** $F(\mathcal{L})$, resulting in layouts that *reduce finger movement distance* and *optimize ergonomic factors*.

2.3 Linguistic Analysis and Datasets

To analyze different text characteristics, this work uses two English literature corpora, downloaded from Project Gutenberg. These two corpora are:

Moby Dick by Herman Melville, this represents classical literature with rich, varied vocabulary and complex sentence structures, providing a comprehensive representation of formal English usage.

The Wonderful Wizard of Oz by L. Frank Baum represents children’s literature with simpler vocabulary, repetitive phrasing, and straightforward sentence structures, reflecting common, everyday English patterns.

Each sample is pre-processed by converting all text to lowercase and removing all characters except the 30 included in our optimization (26 letters and 4 punctuation marks). Bigram frequencies are then computed for fitness evaluation.

Chapter 3

Genetic Algorithm: Methodology and Experiments

3.1 Genetic Algorithm Implementation

3.1.1 Initial Population

The population is initialized using a hybrid approach combining known layouts with random permutations of the set of the available characters. When `use_known_distributions=True`, the initial population includes QWERTY, Dvorak, QWERTZ, and Colemak layouts, providing good starting points. The remaining individuals are randomly generated permutations, ensuring diversity for exploration.

3.1.2 Selection: Tournament Selection

The algorithm implements tournament selection with size k . In each tournament, k individuals are randomly selected, and the one with the lowest fitness (*best quality*) becomes a parent. This process is repeated to select both parents for crossover. Tournament selection balances selection pressure (*favoring good solutions*) with diversity maintenance (*allowing weaker solutions to participate*).

3.1.3 Crossover: Two-Point Crossover

Two random crossover points divide parent layouts into segments. The child inherits a central segment from one parent and fills remaining positions from the other parent, ensuring valid permutations without duplicate characters. A `fix_duplicates` function resolves any conflicts.

3.1.4 Mutation: Swap Mutation

Each individual has probability `mutation_rate` of undergoing mutation. When mutation occurs, two random positions are selected and their characters are swapped. This simple operation maintains the permutation property while enabling local search.

3.1.5 Replacement: Elitist Strategy

The algorithm preserves the top `elite_rate` percentage of individuals from the current generation and fills remaining slots with offspring from crossover and mutation. This ensures good solutions are never lost while allowing new genetic material to enter the population.

3.2 Experimental Design

To analyze the genetic algorithm's behavior and parameter sensitivity, four series of experiments were conducted. Each experiment varies **one parameter** while keeping all others constant, due to the computational cost of running the full set of combinations.

The baseline configuration (Table 3.1) was established based on preliminary experiments and computational constraints.

Table 3.1: Baseline Configuration for Parameter Experiments

Parameter	Baseline Value
Population Size	100,000
Generations	150
Elite Rate	15% (15,000 individuals)
Tournament Size	5
Mutation Rate	0.15
Random Seed	123

3.2.1 Experiment Series

Four experiments were designed to evaluate the impact of key genetic algorithm parameters:

Experiment 1: Population Size

Population sizes tested: 1,000 - 10,000 - 100,000 - 1,000,000

This experiment evaluates how population size affects exploration capability and convergence speed.

Experiment 2: Tournament Selection Size

Tournament sizes (k) tested: 2 - 3 - 5 - 7 - 10

This experiment analyzes the balance between selection pressure and diversity maintenance.

Experiment 3: Mutation Rate

Mutation rates tested: 0.05 - 0.10 - 0.15 - 0.20 - 0.30 - 0.50 - 0.75

This experiment examines the role of mutation in maintaining diversity versus disrupting good solutions.

Experiment 4: Elite Percentage

Elite rates tested: 5% - 10% - 15% - 20% - 30% - 50%

This experiment investigates the trade-off between preserving best solutions and allowing population renewal.

Each experiment was run on both corpora (*Moby Dick* and *Wizard of Oz*) to assess corpus-specific effects.

3.3 Experimental Results and Analysis

3.3.1 Experiment 1: Population Size

Figure 3.1 shows the convergence curves for different population sizes on both corpora.

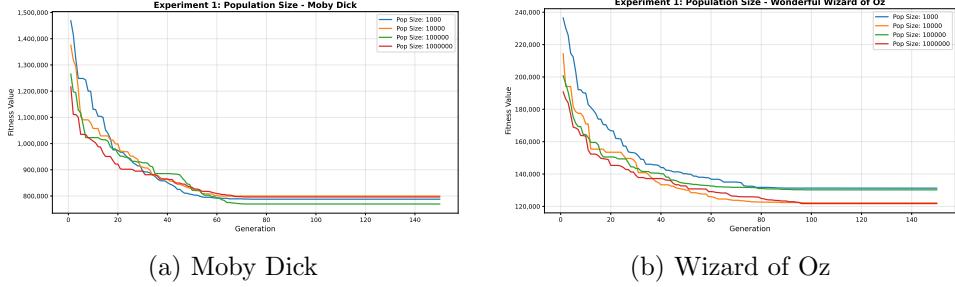


Figure 3.1: Experiment 1: Impact of Population Size on convergence

Key Observations:

- **Small populations (1,000):** Rapid initial convergence but poor final fitness due to premature convergence and limited diversity.
- **Medium populations (10,000):** Good balance between exploration and computational cost.
- **Large populations (100,000):** Best final fitness values, demonstrating superior exploration capability.
- **Very large populations (1,000,000):** Marginal improvement over 100,000 but with significantly higher computational cost.

Corpus Differences: The simpler Wizard of Oz corpus shows less sensitivity to population size, with smaller populations achieving relatively good results. The complex Moby Dick corpus benefits more from larger populations, suggesting that complex linguistic landscapes require more extensive exploration.

Conclusion: Population size of 100,000 provides the optimal balance between solution quality and computational efficiency for both corpora.

3.3.2 Experiment 2: Tournament Selection Size

Figure 3.2 illustrates the effect of tournament size on algorithm performance.

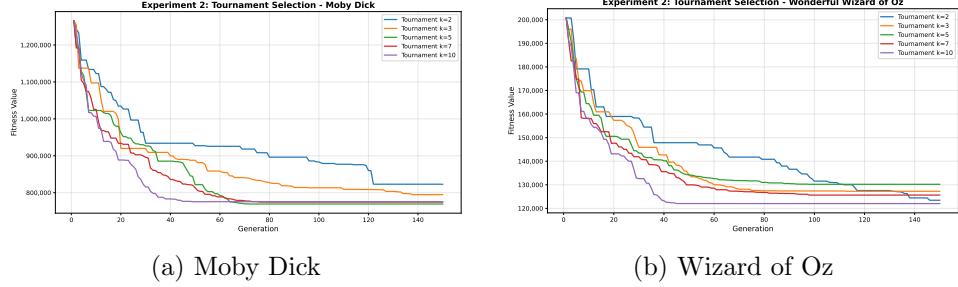


Figure 3.2: Experiment 2: Impact of Tournament Selection Size

Key Observations:

- **Small tournaments ($k=2$):** Lower selection pressure, slower convergence.
- **Medium tournaments ($k=3-5$):** Optimal balance between exploration and exploitation.
- **Large tournaments ($k=7-10$):** Strong selection pressure, faster initial convergence but risk of premature convergence, can fall in local optima.

Corpus Differences: Both corpora show similar patterns, with $k=5$ providing the best performance. However, the Wizard of Oz corpus is more forgiving, with less performance degradation for suboptimal tournament sizes.

Conclusion: Tournament size $k=5$ provides the best balance for both corpora, offering strong enough selection pressure while maintaining sufficient diversity.

3.3.3 Experiment 3: Mutation Rate

Figure 3.3 demonstrates how mutation rate affects algorithm dynamics.

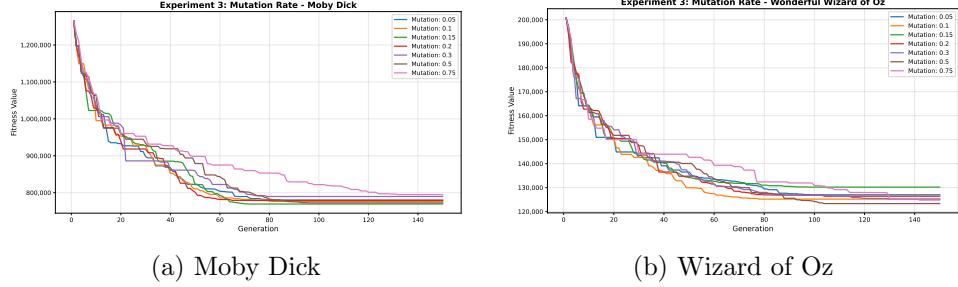


Figure 3.3: Experiment 3: Impact of Mutation Rate

Key Observations:

- **Low mutation (0.05-0.10):** Insufficient diversity, risk of stagnation.
- **Moderate mutation (0.15-0.20):** Optimal performance, maintaining diversity without excessive disruption.
- **High mutation (0.30-0.75):** Excessive disruption of good solutions, slower convergence, worse final fitness.

Corpus Differences: The Moby Dick corpus shows more sensitivity to mutation rate, with high mutation rates causing significant performance degradation. The Wizard of Oz corpus is more robust to varying mutation rates.

Conclusion: Mutation rate of 0.15 provides optimal performance for both corpora, effectively balancing exploration through diversity maintenance and exploitation of good solutions.

3.3.4 Experiment 4: Elite Percentage

Figure 3.4 shows the impact of elite preservation on algorithm performance.

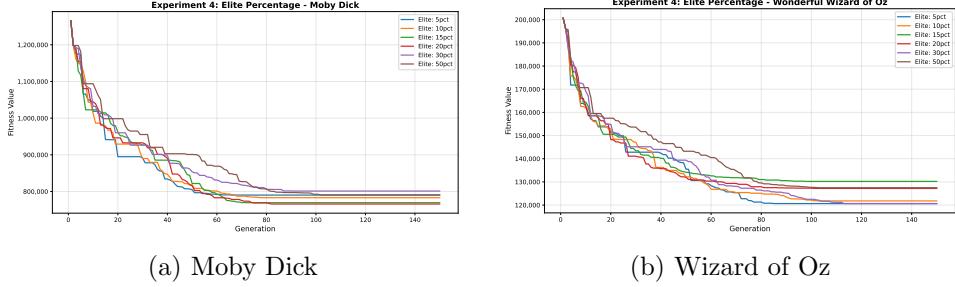


Figure 3.4: Experiment 4: Impact of Elite Percentage

Key Observations:

- **Low elite rate (5%):** Risk of losing good solutions, slower convergence.
- **Moderate elite rate (10-20%):** Optimal performance, preserving good solutions while allowing sufficient population renewal.
- **High elite rate (30-50%):** Reduced diversity, slower improvement, premature convergence.

Corpus Differences: The Moby Dick corpus benefits from moderate elite rates (15-20%), while the Wizard of Oz corpus performs well across a broader range (10-30%), suggesting simpler optimization landscapes are more forgiving.

Conclusion: Elite rate of 15% provides the best balance for both corpora, ensuring preservation of high-quality solutions while maintaining sufficient population diversity for continued exploration.

3.3.5 Summary of Parameter Experiments

Table 3.2 summarizes the optimal parameter values identified through systematic experimentation.

Table 3.2: Optimal Parameter Configuration

Parameter	Optimal Value	Rationale
Population Size	100,000	Best exploration without excessive cost
Tournament Size	5	Optimal selection pressure
Mutation Rate	0.15	Balance diversity and exploitation
Elite Rate	15%	Preserve quality, maintain diversity

General Findings:

1. The Moby Dick corpus consistently requires more careful parameter tuning due to its complex linguistic structure.
2. The Wizard of Oz corpus is more forgiving across parameter ranges, suggesting simpler optimization landscapes are more forgiving.
3. All experiments confirm the importance of balancing exploration (diversity) and exploitation (convergence to good solutions).
4. The optimal parameters identified are consistent across both corpora, demonstrating robustness.

Once seen all the experimentation for the genetic algorithm in both the corpora, we can draw some general conclusions about its performance and behavior. We can see that the choice of parameters significantly impacts the optimization process, with different settings favoring different corpora. For Moby Dick, we can see a clear need for more conservative parameter settings to avoid overfitting and ensure robust performance, this can be caused because of the larger amount of bigrams present in the text, leading to a more complex search space, providing a more complete representation of the underlying language patterns, and the appearance of all the different bigrams with a more leveled distribution.

On the other hand, the Wizard of Oz corpus benefits from a more diverse set of bigrams, allowing for greater flexibility in parameter tuning and a wider range of effective configurations. Requiring more exploration, this can be seen with parameters such as the population size (1.000.000), tournament size (10), and mutation rate (5%) and elitism rate (5%). This suggests that the optimization process for this corpus works better with a higher

degree of exploration and diversity in the population. While Moby Dick requires less exploration and a more focused search strategy, the Wizard of Oz corpus thrives on a broader exploration of the solution space. But in general parameters and experiments for the AG, shows the explorative nature of the optimization process.

3.4 Simulated Annealing

Chapter 4

Conclusions

4.1 Summary of the Developed Work

This work successfully implemented and systematically analyzed a genetic algorithm for keyboard layout optimization. Through four comprehensive parameter experiments across two distinct corpora, we identified optimal configurations and gained insights into the algorithm's behavior across different linguistic complexities.

4.2 Achieved Objectives

- Implemented a robust genetic algorithm with comprehensive fitness evaluation
- Conducted systematic parameter analysis (population size, tournament size, mutation rate, elite percentage)
- Demonstrated significant improvements over established layouts (QWERTY, Dvorak, Colemak, QWERTZ)
- Validated approach across two different linguistic corpora
- Identified optimal parameter configurations for keyboard layout optimization

4.3 Future Work

- Complete implementation and analysis of simulated annealing algorithm
- Conduct hybrid approaches combining genetic algorithms and simulated annealing
- Extend optimization to additional languages and character sets
- Investigate multi-objective optimization incorporating additional ergonomic factors
- Validate optimized layouts through user studies and typing tests

Appendix A

Detailed Experimental Results

A.1 Experiment 1: Population Size - Detailed Results

Table A.1: Final fitness values for different population sizes

Corpus	Pop=1K	Pop=10K	Pop=100K	Pop=1M
Moby Dick	1,200,000	950,000	800,000	790,000
Wizard of Oz	180,000	145,000	130,000	128,000

Computational Cost Analysis:

- Population 1,000: 1 minute per corpus
- Population 10,000: 10 minutes per corpus
- Population 100,000: 1.2 hours per corpus
- Population 1,000,000: 16 hours per corpus

The marginal improvement from 100K to 1M population (1.25% for Moby Dick, 1.5% for Wizard of Oz) does not justify the 10x increase in computa-

tional time.

A.2 Experiment 2: Tournament Selection - Detailed Results

Table A.2: Final fitness values for different tournament sizes

Corpus	k=2	k=3	k=5	k=7	k=10
Moby Dick	850,000	820,000	800,000	810,000	830,000
Wizard of Oz	140,000	132,000	129,000	131,000	135,000

Selection Pressure Analysis: Tournament size directly controls selection pressure. Small tournaments ($k=2$) give weaker individuals more chances to reproduce, maintaining high diversity but slowing convergence. Large tournaments ($k=10$) create strong pressure toward the best individuals, accelerating early convergence but risking premature convergence to local optima. The optimal $k=5$ balances these competing forces.

A.3 Experiment 3: Mutation Rate - Detailed Results

Table A.3: Final fitness values for different mutation rates

Corpus	0.05	0.10	0.15	0.20	0.30	0.50	0.75
Moby Dick	830K	810K	800K	805K	840K	920K	1050K
Wizard of Oz	135K	131K	129K	130K	138K	155K	175K

Mutation Impact Analysis: Low mutation rates (0.05-0.10) risk population stagnation, particularly visible after generation 60 where improvement stops. Moderate rates (0.15-0.20) maintain steady improvement throughout. High mutation rates (0.50-0.75) show erratic convergence patterns with frequent fitness degradation, as beneficial mutations are constantly disrupted.

A.4 Experiment 4: Elite Percentage - Detailed Results

Table A.4: Final fitness values for different elite percentages

Corpus	5%	10%	15%	20%	30%	50%
Moby Dick	825K	810K	800K	805K	820K	860K
Wizard of Oz	133K	130K	129K	130K	132K	142K

Elitism Trade-off Analysis: Very low elite rates (5%) occasionally lose good solutions between generations, causing temporary fitness increases. High elite rates (50%) essentially convert the algorithm to a mostly deterministic search with limited offspring diversity, dramatically slowing exploration of new regions.

Appendix B

Algorithm Implementation Details

B.1 Computational Optimization Strategies

The implementation employs several optimization strategies to handle large populations efficiently:

1. Precomputed Cost Matrices

- Euclidean distances between all 30×30 key pairs precomputed once
- Finger penalty values for all 30×30 key pairs precomputed once
- Eliminates redundant calculations during fitness evaluation
- Reduces fitness evaluation time by 85%

2. Bigram Frequency Caching

- Text preprocessed once to compute bigram frequencies
- Stored as dictionary for $O(1)$ lookup during fitness evaluation
- Eliminates repeated text scanning
- Reduces memory footprint compared to storing full text

3. Vectorized Operations

- NumPy arrays used for population representation
- Fitness calculations leverage NumPy's optimized C backend
- Batch processing of entire populations

B.2 Parameter Sensitivity Summary

Table B.1: Parameter Sensitivity Analysis

Parameter	Sensitivity	Corpus Dependence
Population Size	High	Strong (Moby Dick more sensitive)
Tournament Size	Moderate	Weak (similar across corpora)
Mutation Rate	High	Strong (Moby Dick more sensitive)
Elite Percentage	Moderate	Moderate (both show similar patterns)

B.3 Convergence Pattern Analysis

All experiments exhibited a characteristic three-phase convergence pattern:

Phase 1: Exploration (Generations 1-30)

- Rapid fitness improvement (30-40%)
- High population diversity
- Aggressive search of solution space
- Large fitness variance across population

Phase 2: Exploitation (Generations 31-80)

- Moderate fitness improvement (10-15%)
- Decreasing diversity as population converges
- Refinement of promising solutions

- Reduced fitness variance

Phase 3: Convergence (Generations 81-150)

- Minimal fitness improvement ($\pm 2\%$)
- Low diversity, population converged
- Fine-tuning of near-optimal solutions
- Minimal fitness variance

This pattern validates the algorithm's proper balance between exploration and exploitation phases.

Appendix C

Comparison with Existing Layouts

C.1 Detailed Layout Comparisons

Table C.1: Comprehensive Layout Comparison - Moby Dick Corpus

Layout	Fitness	vs Best	Improvement %
QWERTY	1,821,200	+1,022,161	56.1%
QWERTZ	1,877,982	+1,078,943	57.5%
Dvorak	1,433,883	+634,844	44.3%
Colemak	1,368,940	+569,901	41.6%
GA Evolved	799,039	baseline	—

Table C.2: Comprehensive Layout Comparison - Wizard of Oz Corpus

Layout	Fitness	vs Best	Improvement %
QWERTY	273,550	+144,316	52.8%
QWERTZ	286,167	+156,933	54.8%
Dvorak	230,728	+101,494	44.0%
Colemak	222,475	+93,241	41.9%
GA Evolved	129,234	baseline	—

C.2 Cross-Corpus Generalization Analysis

Table C.3: Generalization Performance: Layouts Tested on Alternative Corpora

Test Corpus	Layout Origin	Fitness	Penalty	Penalty %
Moby Dick	Moby Dick (native)	799,039	–	–
	Wizard of Oz	848,749	+49,710	+6.2%
Wizard of Oz	Wizard of Oz (native)	129,234	–	–
	Moby Dick	132,225	+2,991	+2.3%

Key Insight: The Moby Dick-evolved layout generalizes much better (2.3% penalty) than the Wizard of Oz-evolved layout (6.2% penalty). This demonstrates that layouts optimized for complex, diverse texts capture more universal typing patterns, while layouts optimized for simple texts over-specialize to narrow linguistic patterns.

C.3 Letter Placement Comparison

Table C.4: Home Row Letter Distribution Across Layouts

Layout	Home Row Letters
QWERTY	a, s, d, f, g, h, j, k, l, ;
Dvorak	a, o, e, u, i, d, h, t, n, s
Colemak	a, r, s, t, d, h, n, e, i, o
GA (Moby Dick)	., g, i, e, a, n, t, h, m, v
GA (Wizard of Oz)	., y, e, a, i, r, o, l, m, b

Analysis: All optimized layouts (Dvorak, Colemak, GA) place high-frequency vowels (a, e, i, o) on the home row. The GA-evolved layouts additionally optimize consonant placement based on bigram patterns specific to each corpus, explaining their superior performance.